

Leo Yao

📍 Palo Alto, CA 📩 leoyao@cmu.edu ☎ +1 (650) 285-0101 ⚙️ leowo 🌐 LeoY20

Education

Carnegie Mellon University <i>B.S. in Statistics & Machine Learning and Computer Science</i>	August 2024 – Present Pittsburgh, PA
<ul style="list-style-type: none">◦ QPA: 3.68/4.00◦ Relevant Coursework: Introduction to Computer Systems, Introduction to Machine Learning, Principles of Functional Programming, Introduction to Computer Security, Principles of Imperative Computation, Probability and Statistical Inference I/II, Statistical Computing◦ Awards: Dean's List, High Honors (Spring 2025, Fall 2025)	

Experience

Research Assistant <i>FastML Lab</i>	September 2025 - Present Pittsburgh, PA
<ul style="list-style-type: none">◦ Implemented a high-granularity quantization-aware training pipeline for FPGA-based GNNs used by CERN's CMS experiment with PyTorch, PyG, and HGQ2 while optimizing for microsecond-latency performance.◦ Developing transformers for real-time rapid reconstruction of High-Granularity Calorimeter energy clusters.	

Machine Learning Research Assistant <i>CMU SPICE Lab</i>	June 2025 – August 2025 Pittsburgh, PA
<ul style="list-style-type: none">◦ Engineered a data preprocessing pipeline using R Tidyverse, Pandas, NumPy, and SciKit-Learn to clean and normalize large-scale energy consumption datasets.◦ Developed a neural network with PyTorch to forecast households' annual cooling energy, engineering a 75% accuracy improvement by implementing LightGBM-based feature selection and hyperparameter tuning (e.g. learning rate annealing, early stopping).	

Projects

Cache Simulator	June 2025
<ul style="list-style-type: none">◦ Developed a write-back, write-allocate cache simulator in C with LRU eviction policy.◦ Designed a flexible architecture to parse instruction traces and accurately model hit/miss/eviction rates across various cache geometries.	
Dynamic Memory Allocator	July 2025
<ul style="list-style-type: none">◦ Developed a dynamic memory allocator using segregated free lists in C. Achieved 74.3% memory utilization and 15885 kilo-operations per second, ranking first in the class by performance.	
Tiny Linux Shell	July 2025
<ul style="list-style-type: none">◦ Built a lightweight Linux shell in C featuring job control and I/O redirection.◦ Improved performance and mitigated race conditions by implementing signal masks and custom signal handlers to manage background process synchronization.	
Web Proxy Server	July 2025
<ul style="list-style-type: none">◦ Created a web proxy server in C. Optimized performance and reduced latency by introducing concurrency with POSIX Threads to handle individual client HTTP GET requests.	
Shark File System	July 2025
<ul style="list-style-type: none">◦ Developed a thread-safe concurrent FAT-style file system in C with file creation, deletion, renaming, reading, and writing capabilities.◦ Optimized performance and correctness with mutex synchronization to mitigate race conditions.	

Skills

Programming Languages: Python, C, SQL, Java, SML, Dafny, x86-64 Assembly, R
Frameworks/Libraries: Pandas, NumPy, Scikit-Learn, PyTorch, Matplotlib, Tidyverse
Developer Tools: Git, Linux, PortSwigger Burp Suite, GDB, Valgrind