

# Leo Yao

📍 Palo Alto, CA 📩 leoyao@cmu.edu ☎ +1 (650) 285-0101 💬 LeoY20

## Education

<b>Carnegie Mellon University</b> <i>B.S. in Statistics &amp; Machine Learning and Computer Science</i>	August 2024 – Present Pittsburgh, PA
<ul style="list-style-type: none"><li>◦ <b>QPA:</b> 3.64/4.00</li><li>◦ <b>Relevant Coursework:</b> Principles of Imperative Computation, Introduction to Computer Systems, Principles of Functional Programming, Introduction to Computer Security, Matrices and Linear Transformations, Calculus in 3D, Probability and Statistical Inference I, Statistical Graphics and Visualization.</li><li>◦ <b>Awards:</b> Dean's List, High Honors (Spring 2025)</li></ul>	

## Experience

<b>Research Assistant</b> <i>FastML Lab</i>	September 2025 - Present Pittsburgh, PA
<ul style="list-style-type: none"><li>◦ Researching deployment of real-time machine learning algorithms for the CMS experiment at CERN LHC.</li><li>◦ Developing high-granularity quantization-aware training features for <a href="#">flowGNN</a>, a library for deploying various types of Graph Neural Networks on FPGA chips in real-time.</li></ul>	

  

<b>Machine Learning Researcher</b> <i>CMU SPICE Lab</i>	June 2025 – August 2025 Pittsburgh, PA
<ul style="list-style-type: none"><li>◦ Cleaned and preprocessed household energy use data with R Tidyverse, Pandas, NumPy, and SciKit-Learn.</li><li>◦ Developed an artificial neural network (ANN) in PyTorch to predict households' annual cooling energy, engineering a performance improvement of over 75% by implementing LightGBM-based feature selection and advanced hyperparameter tuning (e.g. learning rate annealing, early stopping).</li><li>◦ Visualized model performance using Matplotlib; presented graphics in a poster to peers.</li><li>◦ Managed codebase and tracked experiments using Git. Poster and code available upon request.</li></ul>	

## Projects

<b>Computer Systems (in C language)</b>	May 2025 – July 2025
<ul style="list-style-type: none"><li>◦ <b>Dynamic Memory Allocator:</b> Developed a dynamic memory allocator in C for Linux. Achieved 74.3% memory utilization and a throughput of 15885 kilo-operations per second, ranking at the top of the class.</li><li>◦ <b>Tiny Shell:</b> Built a tiny Linux shell with foreground/background job control and I/O redirection.</li><li>◦ <b>Web Proxy Server:</b> Created a multithreaded web proxy server that facilitates communication between clients and web servers.</li><li>◦ <b>Shark File System:</b> Developed a thread-safe concurrent file system with standard features (e.g. file renaming, file descriptor position management), optimizing multithreaded performance and correctness with mutex synchronization.</li></ul>	

## Leadership

<b>Joint Funding Committee Member</b> <i>CMU Student Government</i>	November 2024 – Present Pittsburgh, PA
<ul style="list-style-type: none"><li>◦ Managing the distribution and allocation of approximately \$2.1 million to 300+ student organizations.</li><li>◦ Acting as a liaison between student organizations and Student Government, advocating for their financial needs during weekly JFC meetings.</li></ul>	

## Skills

<b>Programming Languages:</b> Python, C, R, Java, Dafny, Standard ML, x86-64 Assembly
<b>Frameworks/Libraries:</b> Pandas, NumPy, SciKit-Learn, PyTorch, Matplotlib, Tidyverse, ggplot2
<b>Developer Tools:</b> Git, Linux